

Introduction to the computing@biostat.ucph

Andreas Kryger Jensen

October 2020

Abstract

This is an introduction to the computing facilities at the Section of Biostatistics, Institute of Public Health, University of Copenhagen. It is maintained by Andreas ([aj](mailto:aj@biostat.ucph)) and the latest version can be found at <https://github.com/aejensen/computing-biostat-ucph>.

Contents

1	Introduction	1
2	How to logon	2
3	Using the Slurm scheduler on cox and rao	2
3.1	Useful commands	2
3.2	A full example	3
4	The older servers running without a scheduler	3
5	Getting help	4
6	Miscellaneous	4
6.1	gcc	4
6.2	Installing the rstan package	4
6.3	Faster installation of R packages	5

1 Introduction

The Section of Biostatistics owns two servers for high performance computing. These are named **cox** (official name: `biostatcomp01fl.sund.root.ku.dk`) and **rao** (official name: `biostatcomp02fl.sund.root.ku.dk`) [It's probably best to use the official names when communicating with KU-IT]. The servers are HPE ProLiant DL325 Gen10 Plus acquired in 2020 and each have an AMD EPYC 7702P processor with 64 cores with 2 threads per core and ½ TB of memory. They run Red Hat Enterprise Linux Server release 7.9 (Maipo) and your personal network drive is mounted as `~` at `/home/ifsv/ku-username`.

Usage of cox and rao is based on submitting jobs to a scheduler. The scheduler used is the the Slurm Workload Manager (<https://slurm.schedmd.com/>). The following point is extremely important:

It is **strictly forbidden** to use cox and rao outside of the Slurm scheduler! This means that you are **not** allowed to run R interactively on the servers and fork processes using e.g., `mclapply` or other facilities from the parallel package.

The **only** exception is if you need to install R packages for later used by jobs submitted to the scheduler. You can do that by starting up R and installing the packages as usual.

The reason behind this strict requirement is that every employee of the section should have equal access to the available resources as fairly as possible. If you are more anarchistically inclined, the section has three older servers that can be used in a free-for-all environment. More on these in section 4 of this document.

2 How to logon

In order to logon to the servers you must first be added as a linux user on the system. It is *only* KU-IT that can do that, so if you are not already a user on the system, you should contact KU-IT.

If you already have been granted a linux user, you connect through ssh. You do this from a terminal session on either Windows, OS X and Linux. The username and password is your KU username and password. Example:

```
ComputerKing@hawaii ~ % ssh czv146@cox
Password:
Last login: Fri Oct 16 09:14:27 2020 from 10.34.96.111
Information on usage of the computing cluster and examples of how to submit jobs can be found in
/home/ifsvafd/bsafd/computingcluster and O:\computingcluster. Man pages can be found at.
http://gridscheduler.sourceforge.net/htmlman/manuals.html
bash-4.2$
```

If you are connected to the internet at CSS through the wall, you can connect directly. If you are on eduroam, another WiFi, home, or otherwise outside of CSS you must establish a VPN connection to KU first. See <https://kUNET.ku.dk/medarbejderguide/Sider/It/Fjernadgang-vpn.aspx> for information on how to do this.

3 Using the Slurm scheduler on cox and rao

Did I mention that:

It is strictly forbidden to use cox and rao outside of the Slurm scheduler!

A typical workflow:

1. Write an R program that calculates something a single time and check that it works
2. Run the program multiple times using the Slurm scheduler
3. Collect the results from the individual runs and digest

```
-- myProject
-- doSimulation.R
-- logs
-- output
```

2 directories, 1 file

3.1 Useful commands

The following commands are used to communicate with the scheduler and probably the only ones that you will need.

- **squeue** - view information about jobs located in the Slurm scheduler
- **sbatch** - submit a command to Slurm
- **scancel** - cancel a job running on Slurm

Example for checking the status of the system. Login to either cox and rao and execute **squeue**.

```
bash-4.2$ squeue
      JOBID PARTITION    NAME    USER  ST       TIME  NODES NODELIST(REASON)
     2627_0  standard    doSim   czv146  R        0:24      1 biostatcomp01f1
```

2627_1	standard	doSim	czv146	R	0:24	1 biostatcomp01f1
2627_2	standard	doSim	czv146	R	0:24	1 biostatcomp01f1
2627_3	standard	doSim	czv146	R	0:24	1 biostatcomp01f1
2627_4	standard	doSim	czv146	R	0:24	1 biostatcomp01f1

It is seen that the user czv146 is running 5 jobs on biostatcomp01f1 with the name doSim, and they have been running for 24 seconds.

3.2 A full example

Consider a simple R program where we generate some random data and estimate a parameter

```
x <- rnorm(1000)
mean(x)

number_of_tasks <- as.numeric(Sys.getenv("SLURM_ARRAY_TASK_COUNT"))
task_id <- as.numeric(Sys.getenv("SLURM_ARRAY_TASK_ID"))

# Generate random seeds corresponding to the number of tasks
# and use the one for the current task id
set.seed(123457890)
seeds <- sample(1:10^6, size = number_of_tasks, replace = FALSE)
set.seed(seeds[task_id])

# ----- Your R program goes between these lines -----
x <- rnorm(1000)
result <- mean(x)
# -----

# Save the results in an individual file in the output folder
save(result, file = paste('output/result-', task_id, '.RData', sep = ""))
```

4 The older servers running without a scheduler

The Section of Biostatistics still has three older servers: **fisher**, **gauss** and **borel**. You can connect to these through ssh in a similar manner as when connecting to cox and rao.

- **fisher** is from 2013 and is a HP ProLiant BL685c G7 server with 4 AMD Opteron 6380 processors with a total of 32 cores, 64 threads and ½ TB of memory running openSUSE 13.1 (Bottle)
- **gauss** is from 2011 and is an IBM System x3755 M3 server with 4 AMD Opteron 6172 processors with a total of 48 cores and 62.7 GB of memory running Red Hat Enterprise Linux Server release 7.9 (Maipo)
- **borel** is from 2011 and is an IBM System x3755 M3 server with 4 AMD Opteron 6172 processors with a total of 48 cores and 58.8 GB of memory running Red Hat Enterprise Linux Server release 7.9 (Maipo)

These servers are run **without** a workload manager meaning that you can execute programs in anyway that you want e.g., run R interactively and use the multi-core facilities from the parallel package. This implies that you can also **crash** the servers by either submitting too many jobs simultaneously or by exceeding the available memory. Please do not do that.

Crashing the servers implies getting help from KU-IT to restart them, so please **be careful and observant** of other users already running things on these servers. The program **top** which can be executed in the shell is useful to see the current CPU utilization and memory usage and may help you decide whether it is wise to start additional jobs on these servers given their current usage.

It is important to be aware of that due to the different versions of operating systems and gcc not all R packages will work across all the servers if they were install on a specific server. This seems to be especially

pronounced for packages depending on Rcpp and its derivatives. Currently, cox and rao are compatible as well as gauss and borel. If an error occurs due to this (most likely seen as a segmentation fault), you can uninstall the packages and reinstall while being logged on to the servers that you subsequently wish to use them on.

Note: It is the plan to shut down gauss and borel by the end of the year 2020.

5 Getting help

If you have small problems or need some software or libraries installed on the servers some of us have `sudo` permissions and might be able to help. You can e.g., contact Andreas (aeje) if you can find him. For larger scale problems or if the servers become non-responsive you should contact KU-IT.

It should be noted that **you** are in control of your own R packages on the systems. There are no globally or regularly updated R packages installed on the servers. R packages are installed locally for each user in `~/R` and you manage these yourself.

6 Miscellaneous

6.1 gcc

cox and rao are equipped with gcc version 4.8.5 which is the default version on RHEL 7. To use the newest version of gcc as required when e.g., installing the rstan or compiling stan programs you must execute

```
scl enable devtoolset-9 'bash'
```

This command will spawn a new shell with gcc version 9.3.1 as the default compiler. This is *not* persistent between sessions! You can execute `exit` to return to the previous shell with the old version of gcc.

Example:

```
-bash-4.2$ gcc --version
gcc (GCC) 4.8.5 20150623 (Red Hat 4.8.5-44)
Copyright (C) 2015 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
-bash-4.2$ scl enable devtoolset-9 'bash'
bash-4.2$ gcc --version
gcc (GCC) 9.3.1 20200408 (Red Hat 9.3.1-2)
Copyright (C) 2019 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
bash-4.2$
```

6.2 Installing the rstan package

rstan can be a difficult package to install and the requirements have changed through versions. Currently it works on cox and rao if you set the following flags in your `~/R/Makevars`.

```
CXX14FLAGS=-O3 -march=native -mtune=native -fPIC
CXX14=g++
```

Remember to also enable the 9.3.1 version of gcc before installing rstan. You must also have this enabled in your session when compiling stan scripts.

6.3 Faster installation of R packages

Some R packages require compilation of c/c++ code. You can speed-up this process by using multiple cores. To do this you can add `MAKEFLAGS=-j64` to `~/.R/Makevars` or more generally add `export MAKEFLAGS=-j64` to your `~/.bash_profile`.